

Introduction

Ramin Moazzeni

San Jose State University

CS157a – Database Mgmt Systems

CS 157a: Database Mgmt Systems

Instructor: Ramin Moazzeni, Ramin.Moazeni@sjsu.edu

- Office hours: after class (or by appointment)

Course web site: <https://sjsu.instructure.com>

Texts and readings:

- Ramakrishnan & Gerke, *Database Management Systems*, 3rd ed.
- For SQL: Greenspun, *SQL for Nerds* (<http://philip.greenspun.com/sql/index.html>)
- Other books may be useful:
 - Elmasri, Navathe. “Fundamentals of Database Systems,” Latest Edition
- Oracle References
 - [Oracle SQL Reference](#)
 - [Oracle SQL*Plus User's Guide and Reference](#)
 - [Online Oracle Documentations](#)
- Various papers and readings
- Prerequisites:
 - [CS 146](#) (with a grade of “C-” or better); Computer Science, Applied and Computational Math, Forensic Science: Digital Evidence, or Software Engineering majors only; or instructor consent.

Course Format and Grading

- Quizzes: 10%
- HW Assignments: 20% (Individual)
- Project: 20% (Group)
 - Specification will be provided later
- Midterm exam: 25% (Individual)
- Final exam: 25% (Individual)

Important Dates

- Class starts Tuesday, June 1
Midterm Exam: Tuesday, July 6 (11am-1pm) ONLINE
Class ends Thursday, Aug 5
Final Exam: Thursday, Aug 5 (11am-1pm) ONLINE

Course Outline

(may vary slightly)

Outline

- ER Data model
- The relational data model
- Map ER to Relational model
- Extended ER
- SQL
- OO & ORDBMS
- Application Programming (DB connectivity)
- XML, JSON, XML Schema
- Normalization
- Transactions
- DB Security
- Spatial DBs – If time permits
- NoSQL – If time permits

Cheating/Plagiarism

- Cheating is a serious offense.
- Cheating includes copying on exams or written assignments; obtaining advance copies of exams; outsourcing homeworks or project work; and copying material from the web and including on homeworks without proper attribution.
- You may discuss concepts with your classmates or anyone else, and you are encouraged to do so. However, when it comes to writing the program (even just the 'pseudo-code'), you must do it yourself
- Please see the university's policy regarding academic integrity: <http://www.sjsu.edu/senate/docs/S07-2.pdf>

Outline for Today's Lecture

- Overview of database systems
 - Recommended reading: Introduction of **SQL for Web Nerds**, by Philip Greenspun
<http://philip.greenspun.com/sql/>
 - Read Chapter 1 of the textbook
- Course outline
- What the course is about

What *Is* a Database Management System ?

- Database: collection of files that store the data
 - Entities: students, faculty, courses and classroom
 - Relationships: between entities
 - Students taking courses
 - Faculty teaching courses
- DataBase Management System = DBMS
 - Software that Manages the DB
 - A big program that accesses and updates those files
- Relational DBMS = RDBMS
 - DBMS that is based on a relational model

Where are RDBMS used ?

- Backend for traditional “database” applications
- Backend for large Web sites
- Backend for Web services

Examples of DBMS Usage

- Airlines: reservations and schedules (Expedia)
- Universities: student info, grades
- Banking: customer info and accounts
- Credit Cards: customer info, transactions
- Sales: customer info, inventory (Amazon, EBay)
- Government: taxes, census

Example: Internet Movie Database (IMDB)



Earth's Biggest Movie Database™

[NOW PLAYING](#) [MOVIE / TV NEWS](#) [MY MOVIES](#) [NEW ON DVD](#) [IMDb TV](#) [MESSAGE BOARDS](#) [SHOWTIMES & TICKETS](#)

[Home](#) | [Top Movies](#) | [Photos](#) | [Independent Film](#) | [GameBase](#) | [Browse](#) | [Help](#)

search [more](#) | [tips](#)

[IMDb](#) > [The Big Lebowski \(1998\)](#)



The Big Lebowski (1998)

[★ photos](#) [board](#) [trailer](#) [IMDb Pro details](#)

[Register](#) or [login](#) to rate this title

★★★★★☆☆☆☆☆

User Rating: 8.1/10 ([93,977 votes](#))

[Top 250: #171](#) [more▶](#)

Overview

Director: [Joel Coen](#) [more▶](#)

Writers (WGA): [Ethan Coen](#) (written by) & [Joel Coen](#) (written by)

Release Date: 6 March 1998 (USA) [more▶](#) [view trailer▶](#)

Genre: [Comedy](#) / [Crime](#) / [Mystery](#) [more▶](#)

Tagline: They figured he was a lazy time wasting slacker. They were right. [more▶](#)

Plot Outline: "Dude" Lebowski, mistaken for a millionaire Lebowski, seeks restitution for his bowling buddies to help get it. [more▶](#)

[add to My Movies](#)

Quicklinks

Top Links

- [trailers](#)
- [full cast and crew](#)
- [trivia](#)
- [official sites](#)
- [memorable quotes](#)

Overview
[main details](#)

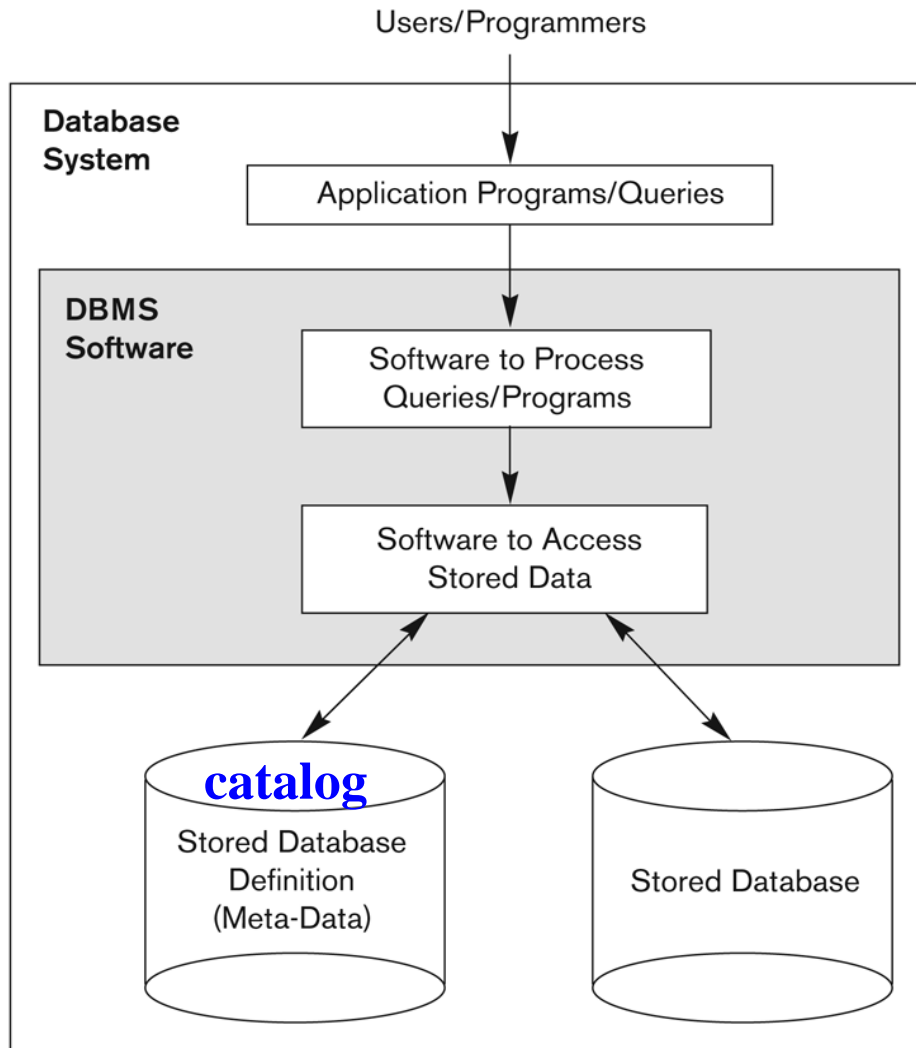
- [combined details](#)
- [full cast and crew](#)
- [company credits](#)

[SHOP](#) 




adver

How is a RDBMS used ?



**Databases are self-describing:
catalog describes the structure
of the data stored in the DB**


Example of a Traditional Database Application

- Suppose we are building a system to store information about:
 - students
 - courses
 - professors
 - who takes what, who teaches what


Can we do it without a DBMS ?

Yes...Start by storing the data in files:


students.txt



courses.txt



professors.txt



Now write C or Java programs to implement specific tasks (i.e store, modify and query)

Doing it without a DBMS...

- Enroll “John Smith” in “COEN444”:

Write a program to do the following:

Read ‘students.txt’

Read ‘courses.txt’

Find&update the record “John Smith”


Find&update the record “COEN444”

Write “students.txt”

Write “courses.txt”

Problems without an DBMS...

- System crashes:



```
Read 'students.txt'  
Read 'courses.txt'  
Find&update the record "John Smith"  
Find&update the record "COEN444"  
Write "students.txt"  
Write "courses.txt"
```

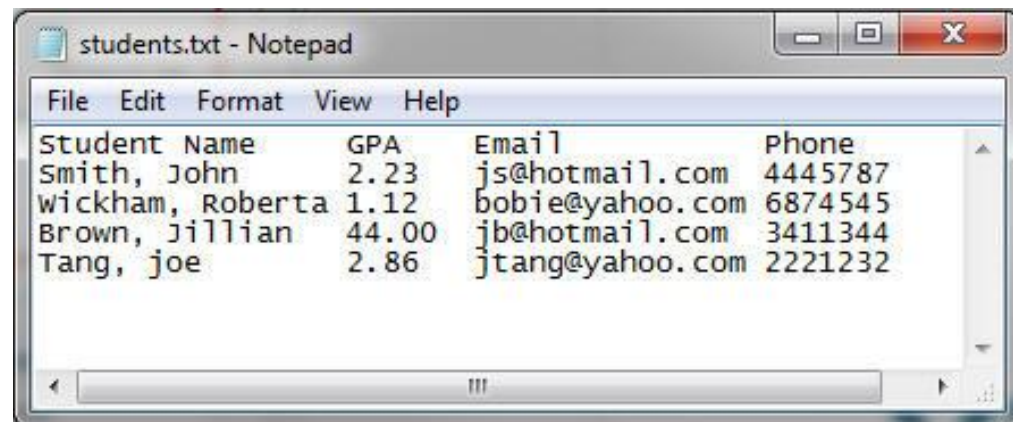
CRASH
!

- What is the problem ?
- Large data sets (say 50TB)
 - What is the problem ?
- Simultaneous access by many users
- Enforcing Constraints
- Scalability
- Security

Enforcing Constraints

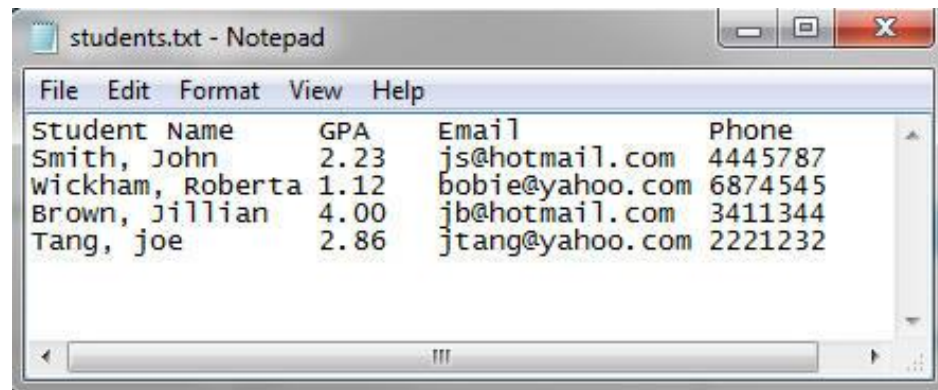
- With the text file solution, there is no way to enforce integrity constraints on the data. In other words people can put bad data into the text file.
- In contrast, a DBMS allows us to enforce all kinds of constraints. This really helps (but does not guarantee) that our data is correct.

A typo gives Jillian Brown a GPA of 44.00



Scalability

- The text file solution might work for small datasets. What happens when we have big datasets?
- Most real world datasets are so large that we can only have a small fraction of them in main memory at any time, the rest has to stay on disk.




A screenshot of a Notepad window titled "students.txt - Notepad". The window displays a table with four columns: Student Name, GPA, Email, and Phone. The data is as follows:

Student Name	GPA	Email	Phone
Smith, John	2.23	js@hotmail.com	4445787
Wickham, Roberta	1.12	bobie@yahoo.com	6874545
Brown, Jillian	4.00	jb@hotmail.com	3411344
Tang, joe	2.86	jtang@yahoo.com	2221232

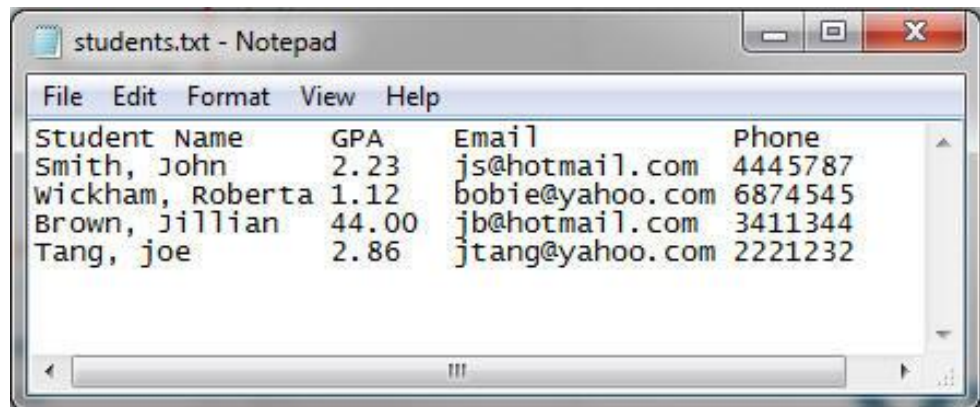
Query Expressiveness

- The text file solution would allow to search for keywords or certain numbers (slowly).
- With a DBMS I can search with much more expressive queries. For example I can ask.. *“Find all students whose GPA is greater than 2.5, and who don’t own a phone”* or *“what is the average GPA of the students”*



Query Expressiveness II

- Could write some program that might allow more expressive queries on the text file, but it would be tied into the structure of my data and the operating system etc..
- With a DBMS we are completely isolated from the physical structure of our data. If we change the structure of our data (by adding a field, for example) or moving from a PC to a Mac, nothing changes at the front end!

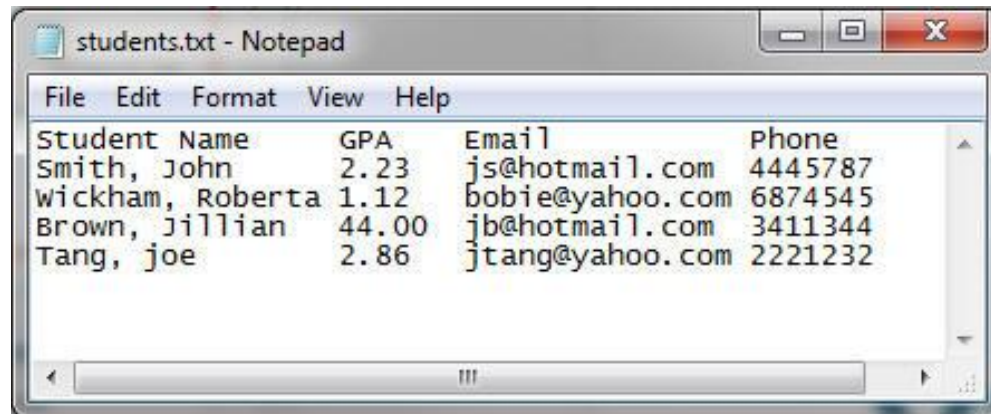


A screenshot of a Notepad window titled "students.txt - Notepad". The window displays a table with four columns: Student Name, GPA, Email, and Phone. The data is as follows:

Student Name	GPA	Email	Phone
Smith, John	2.23	js@hotmail.com	4445787
Wickham, Roberta	1.12	bobie@yahoo.com	6874545
Brown, Jillian	44.00	jb@hotmail.com	3411344
Tang, joe	2.86	jtang@yahoo.com	2221232

Different Views

- The text file solution only allows one view of the data.
- With a DBMS, could arrange for different people to have different views of the data. For example, a professor can see everything, while student can see only his/her data, and a TA can see data for students in his/her section, etc.

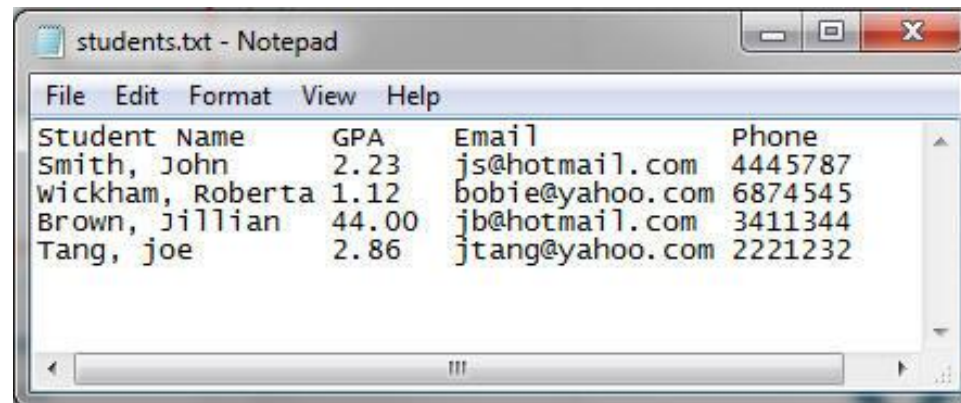


A screenshot of a Notepad window titled "students.txt - Notepad". The window displays a text file with student data organized in a table-like format. The data is as follows:

Student Name	GPA	Email	Phone
Smith, John	2.23	js@hotmail.com	4445787
Wickham, Roberta	1.12	bobie@yahoo.com	6874545
Brown, Jillian	44.00	jb@hotmail.com	3411344
Tang, Joe	2.86	jtang@yahoo.com	2221232

Concurrency

- Suppose the text file is being modified at the same time by multiple users (i.e professor and the TA)
- A DBMS will automatically make sure that this kind of thing cannot happen.

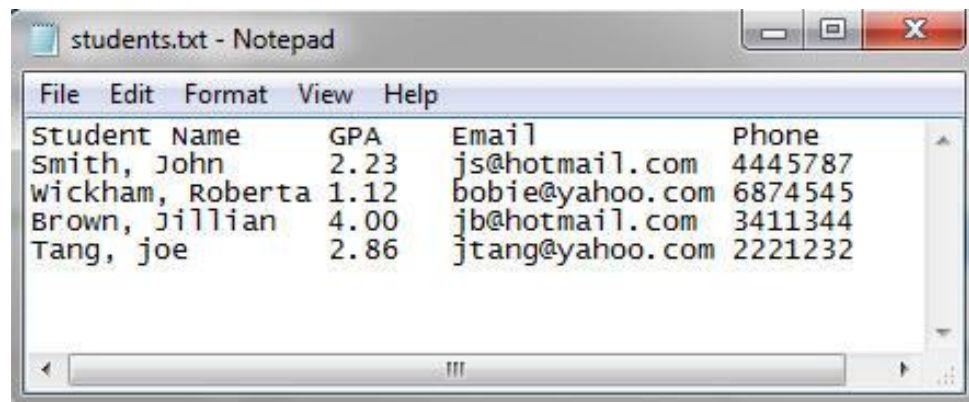


A screenshot of a Notepad window titled "students.txt - Notepad". The window displays a table with four columns: Student Name, GPA, Email, and Phone. The table contains four rows of student data.

Student Name	GPA	Email	Phone
Smith, John	2.23	js@hotmail.com	4445787
Wickham, Roberta	1.12	bobie@yahoo.com	6874545
Brown, Jillian	44.00	jb@hotmail.com	3411344
Tang, joe	2.86	jtang@yahoo.com	2221232

Security

- Suppose the text file on UNIX account, and a student hacks in and changes their grades...
- A DBMS will allow multiple levels of security.
 - Enforce security policies for users to access different subset of data

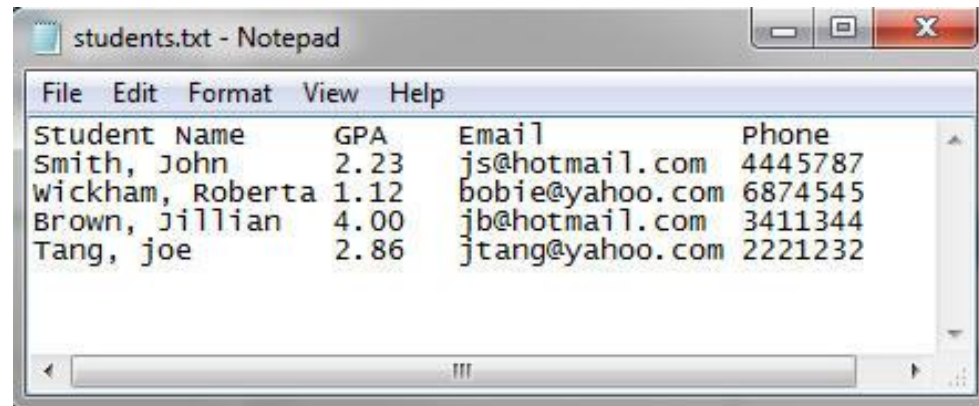


A screenshot of a Notepad window titled "students.txt - Notepad". The window displays a table with four columns: Student Name, GPA, Email, and Phone. The table contains four rows of student data.

Student Name	GPA	Email	Phone
Smith, John	2.23	js@hotmail.com	4445787
Wickham, Roberta	1.12	bobie@yahoo.com	6874545
Brown, Jillian	4.00	jb@hotmail.com	3411344
Tang, joe	2.86	jtang@yahoo.com	2221232


Crash Recovery

- Suppose while modifying the text file, the system crashes!
- A DBMS is able to guarantee 100% recovery from system crashes (to a consistent state).

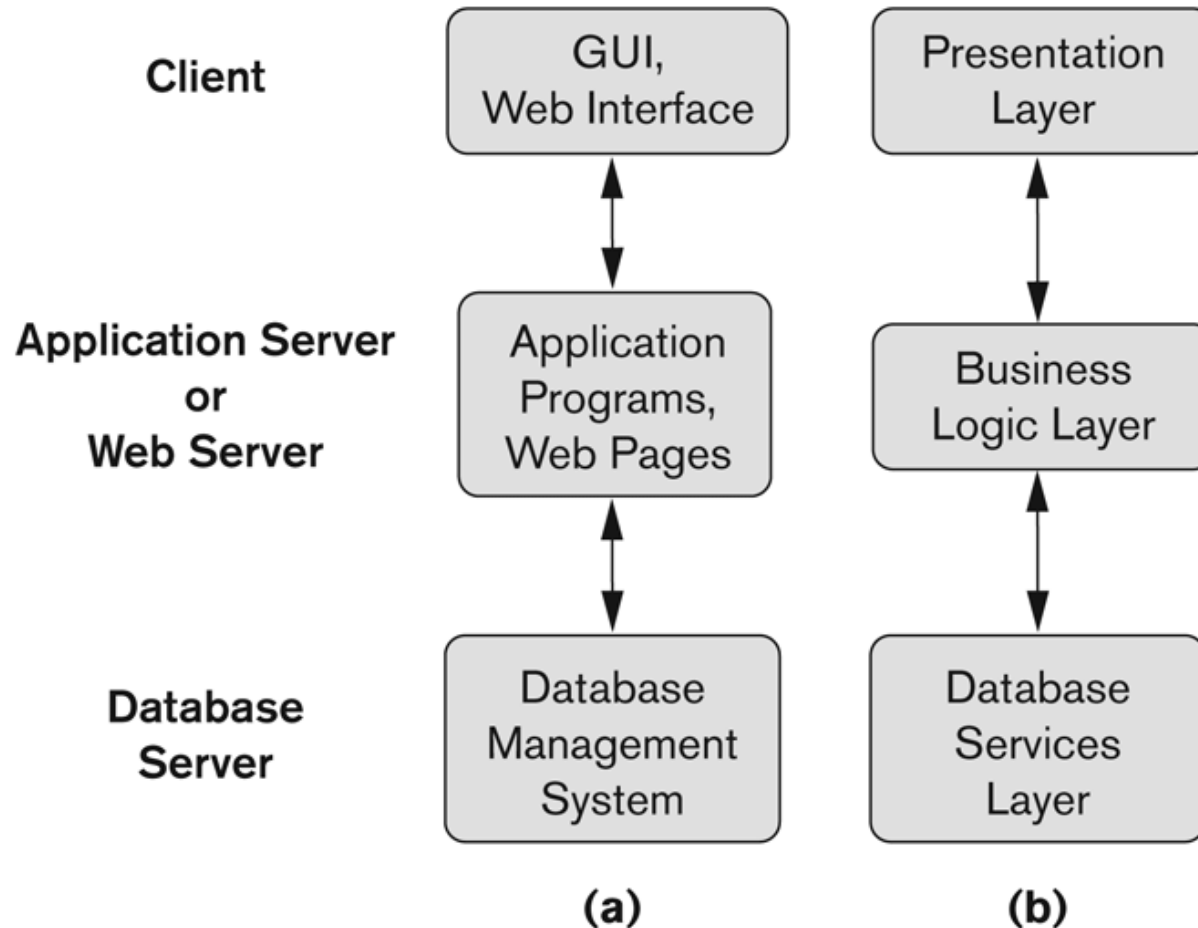


Building the Applications (2-tier)

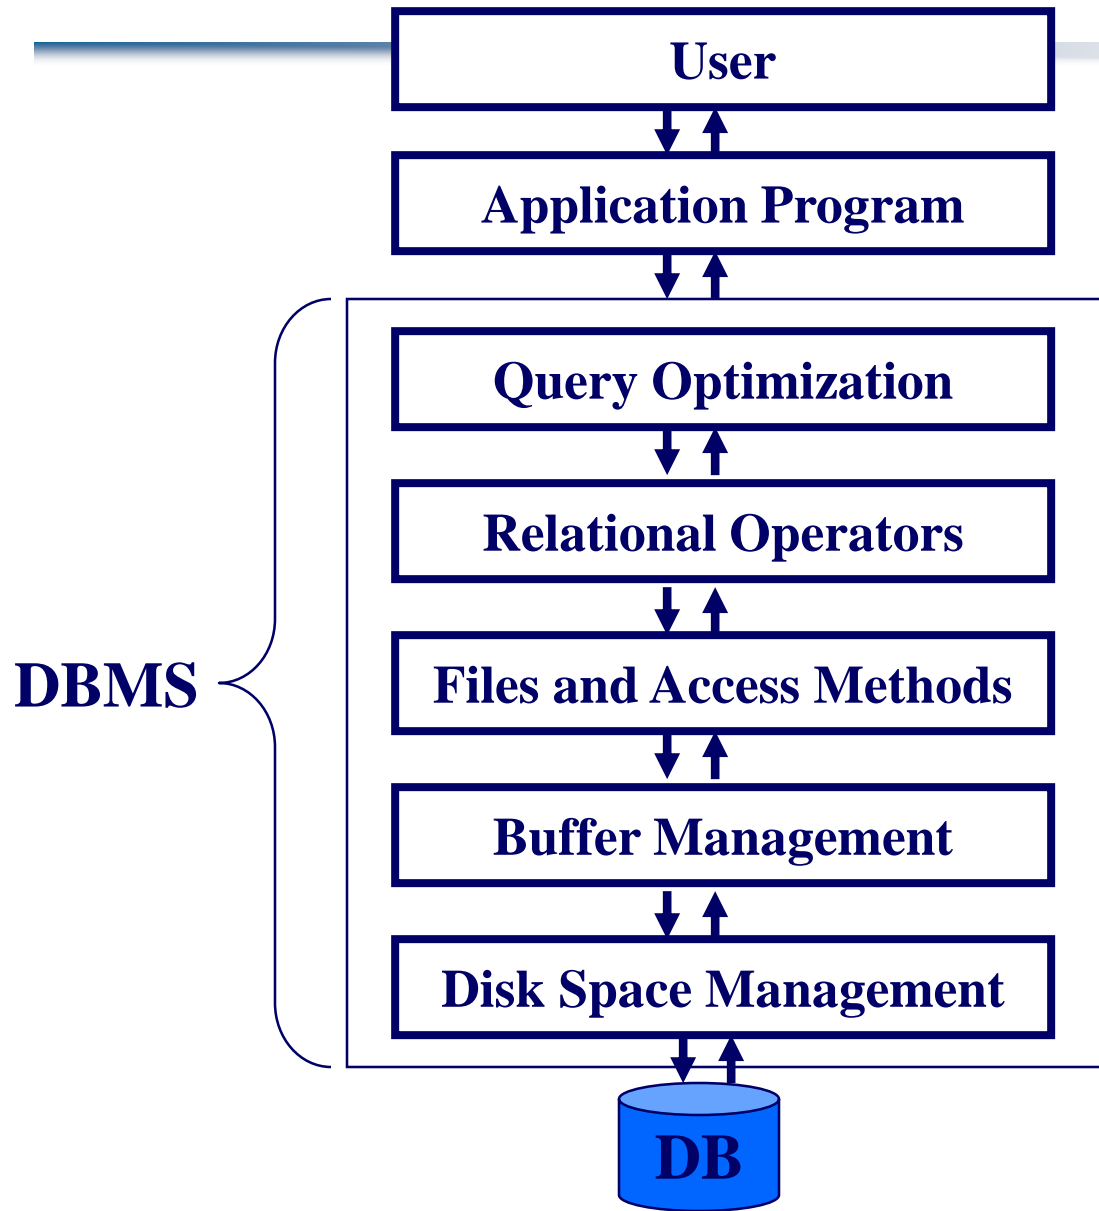
“Two tier database system”



Building the Applications (3-tier)



Control Abstraction



Each layer
need not know
(or care) how
other layers are
implemented

The Database Abstraction Provided by the DBMS

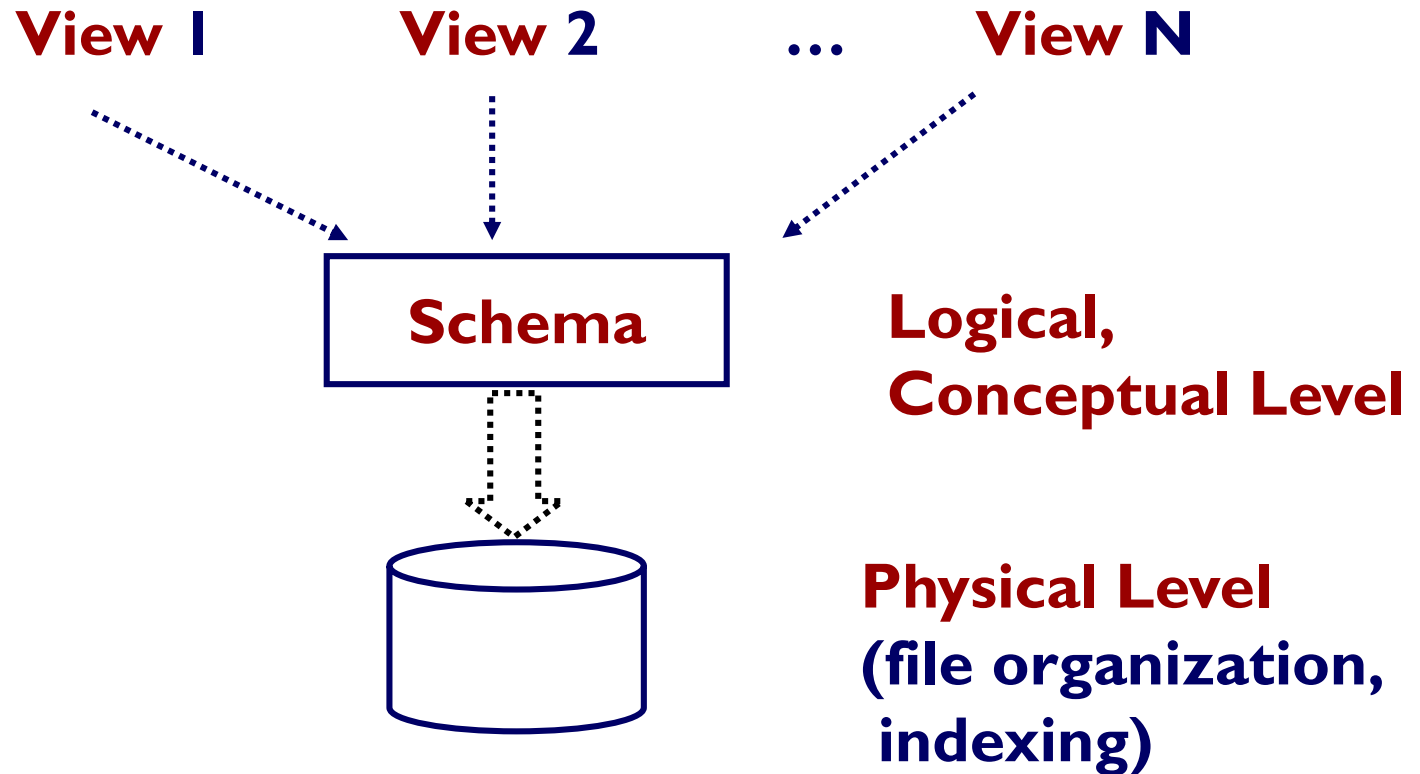
We think of databases at two levels:

- **Logical** structure:
 - What users/programmers see – program or query interface
- **Physical** structure:
 - Organization on disk, indices, etc.

The logical level is further split into:

- Overall database design (conceptual; seen by the DB designer)
- **Views** that various users get to see

The Three-level Architecture for Databases



Functionality of a DBMS

The programmer sees SQL, which has two components:

- Data Definition Language – DDL
 - Creating database tables
- Data Manipulation Language - DML
 - query language

Behind the scenes the DBMS has:

- Query engine
- Query optimizer
- Storage management
- Transaction Management (concurrency, recovery)

How the Programmer Sees the DBMS

- Start with DDL to *create tables*:

```
CREATE TABLE Students (  
    Name CHAR(30)  
    SSN CHAR(9) PRIMARY KEY NOT NULL,  
    Category CHAR(20)  
) ...
```

- Continue with DML to *populate tables*:

```
INSERT INTO Students  
VALUES('Charles', '123456789', 'undergraduate')  
... .
```

How the Programmer Sees the DBMS

- Tables:

Students:

SSN	Name	Category
123-45-6789	Charles	undergrad
234-56-7890	Dan	grad

Takes:

SSN	CID
123-45-6789	CSE444
123-45-6789	CSE541
234-56-7890	CSE142
	...

Courses:

CID	Name	Quarter
CSE444	Databases	fall
CSE541	Operating systems	winter

- Still implemented as files, but behind the scenes can be quite complex

“data independence” = separate logical view from physical implementation

Transactions

- Enroll “John Smith” in “COEN444”:

```
BEGIN TRANSACTION;  
  
INSERT INTO Takes  
  SELECT Students.SSN, Courses.CID  
  FROM Students, Courses  
  WHERE Students.name = ‘John Smith’ and  
         Courses.name = ‘Databases’  
  
-- More updates here....  
  
IF everything-went-OK  
  THEN COMMIT;  
ELSE ROLLBACK
```

If system crashes, the transaction is still either committed or aborted

Transactions

- A *transaction* = sequence of statements that either all succeed, or all fail
- Transactions have the ACID properties:

A = atomicity

each transaction be "all or nothing": if one part of the transaction fails, the entire transaction fails

C = consistency

Any transaction will bring the database from one valid state to another

I = independence

ensures that the concurrent execution of transactions results in a system state that would be obtained if transactions were executed serially

D = durability

once a transaction has been committed, it will remain so, even in the event of power loss, crashes, or errors.

Queries

- Find all courses that “John” takes

```
SELECT C.name  
FROM   Students S, Takes T, Courses C  
WHERE  S.name=“John” and  
        S.ssn = T.ssn and T.cid = C.cid
```

- What happens behind the scene ?
 - Query processor figures out how to answer the query efficiently.

Advantages of a DBMS

- Data Independence

- Logical Data Independence

- Ability to change the logical (conceptual) schema without changing the External schema (User View)

- Physical Data Independence

- Ability to change the physical schema without changing the logical schema
Protection from changes in physical structure of data.

- Query expressiveness

- Reduced application development time.

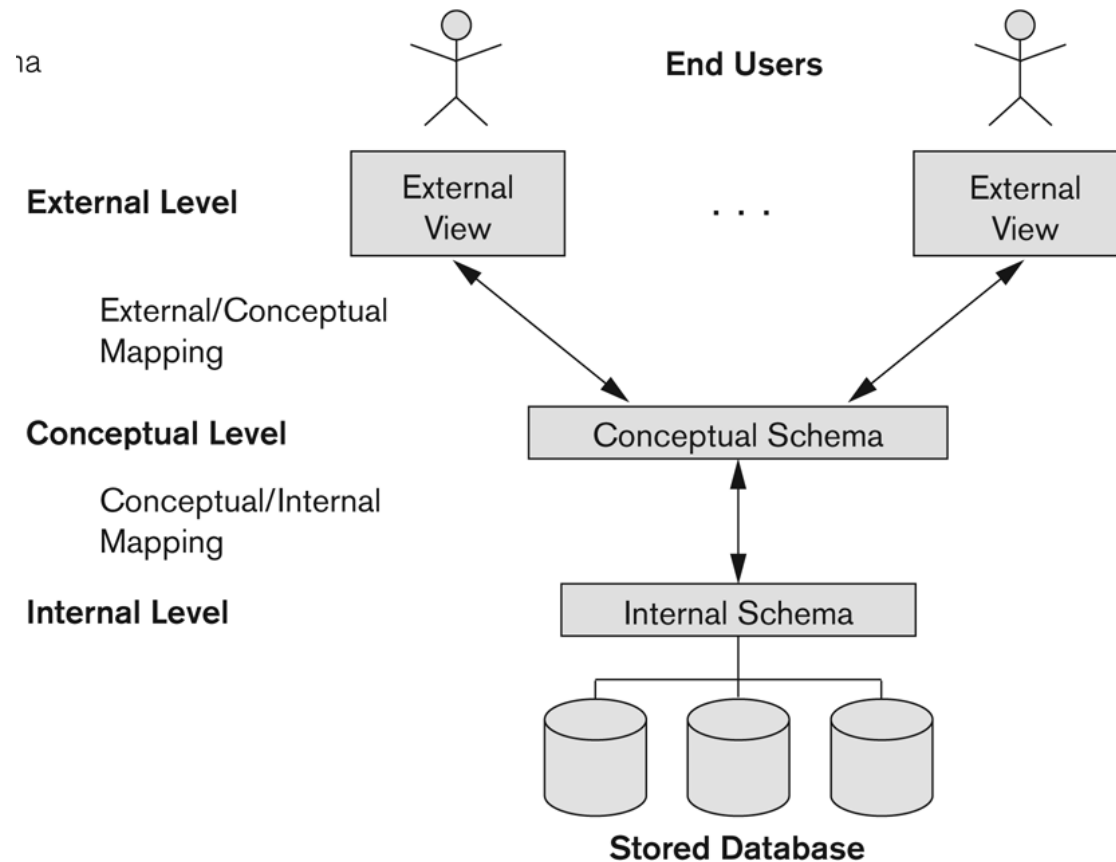
- Concurrency and Crash Recovery

- Schedule concurrent access to the data

- Security

- Enforce access controls

Data Independence



Each layer need not know how other layers organize data

Data Independence

A user of a relational database system should be able to use the database without knowing about how the precisely how data is stored, e.g.

```
SELECT When, Where  
FROM Calendar  
WHERE Who = "Jane"
```

After all, you don't worry about the IEEE floating-point specs when you do division in a Java program or with a calculator

More on Data Independence

Logical data independence

Protects the user from changes in the logical structure of the data:

could reorganize the calendar “schema” without changing how we query it

Physical data independence

Protects the user from changes in the physical structure of data:

could add an index on **who** (or sort by **when**) without changing how the user would write the query, but the query would execute faster (query optimization)

Why Learn about DBMS?

- Many decisions about how to use a DBMS for an application depend on the capabilities of the DBMS
- To use it well, it's necessary to also understand how a DBMS *works*.

Database Users

- End users (or DB application users)
- DB application programmers (more precisely, they are *DBMS* users)
 - E.g. webmasters
- Database administrator (DBA)
 - Designs logical /physical schemas
 - Handles security and authorization
 - Data availability, crash recovery
 - Database tuning as needs evolve

Must understand how a DBMS works!

Data Model

- Data Model:
 - Collection of concepts for describing data
- Schema
 - Description of a particular collection of data, using the a given data model.
- Relational model of data
 - Main concept: relation, basically a table with rows and columns.
 - Every relation has a schema, which describes the columns, or fields

ER Model Basics

- Entity:
 - A real world object or thing
- Attribute
 - Each entity has attributes
- Entity Set
 - Collection of similar entities
 - All entities in an entity set have the same set of attributes

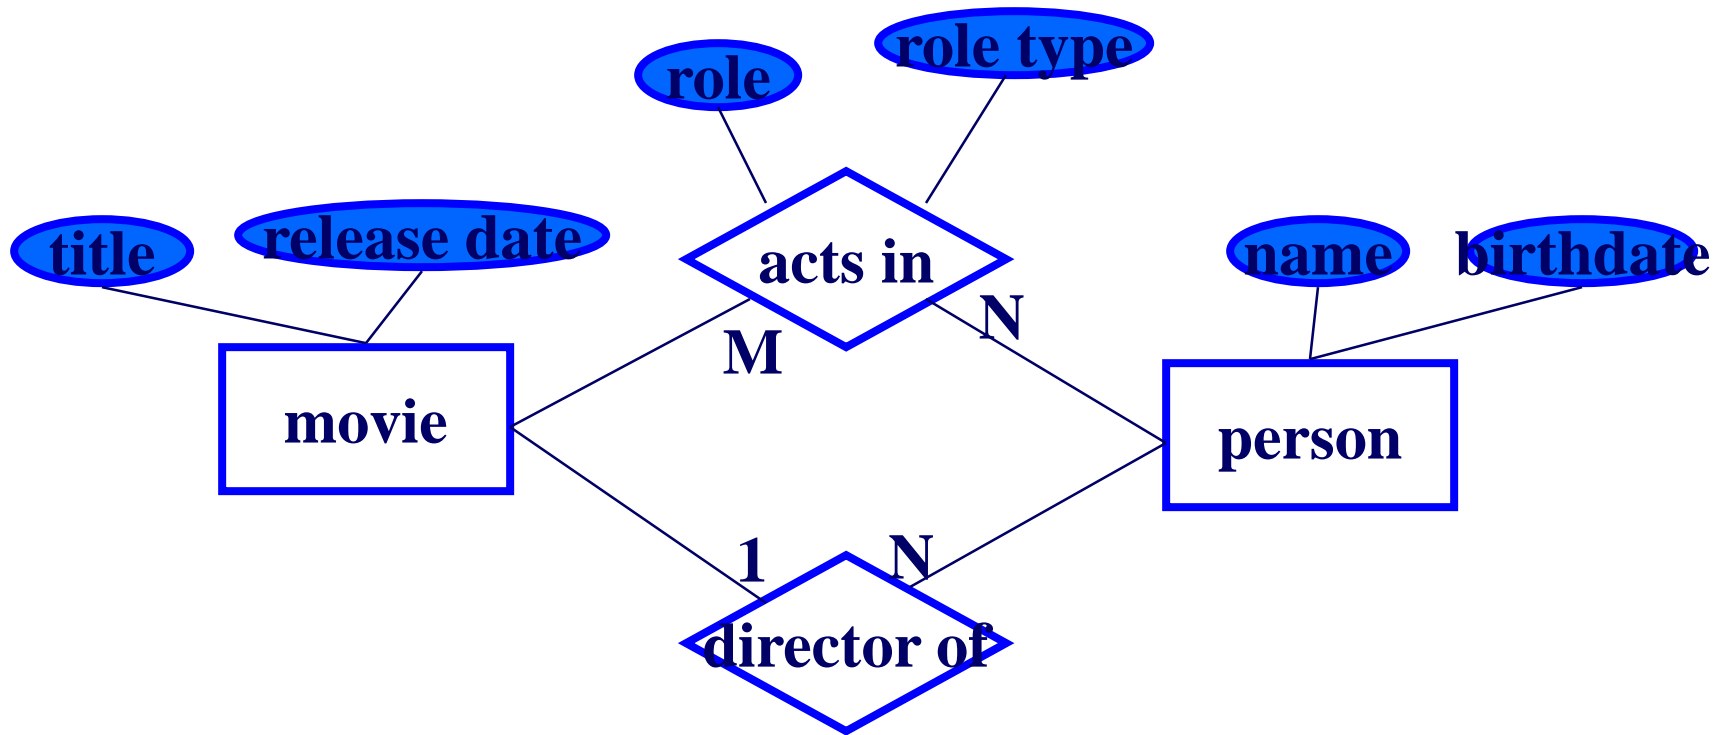
An entity can be uniquely identified thru its attributes

Overview of Database Design

- Conceptual design
 - Use *ER Model*: E- *Entities* and R-*Relationships*
 - Decide the *entities* and *relationships* in the enterprise.
 - Decide what information about these entities and relationships should we store in the database.
 - Decide the *integrity constraints* or *business rules*.
- Implementation (logical design)
 - Map an ER model into a relational schema.

Building a DB: construct a conceptual model

- A conceptual model identifies entities and relationships



entity

attribute

relationship

Building a DB:

Define Relational Schema

- A schema describes DB using *data model* supported by DMBS (eg, *relational model*)
- RDBMS – DBMS that supports relational model

MOVIE

MID	Title	Rating	Director
-----	-------	--------	----------

PERSON

PID	Name	Bday
-----	------	------

ACTS_IN

MID	PID	Role	Rtype
-----	-----	------	-------

Building a DB:

Populate DB

MOVIE

MID	Title	Rating	Director
1	The Big Lebowski	R	72
2	Star Wars	PG	29
...			

ACTS_IN

MID	PID	Role	Rtype
1	1	The Dude	STAR
2	2	Han Solo	CO_STAR
...			

PERSON

PID	Name	Bday
1	Jeff Daniels	12/4/49
2	Harrison Ford	7/13/42
...		

Set initial records of the DB

Querying The Database

- Most RDBMS allow users to query the database using SQL (structured query language)
- Example: get cast of “The Big Lebowski”

```
SELECT Name, Role, Rtype  
FROM PERSON, ACTS_IN  
WHERE MID = '1' AND PERSON.PID == ACTS_IN.PID
```

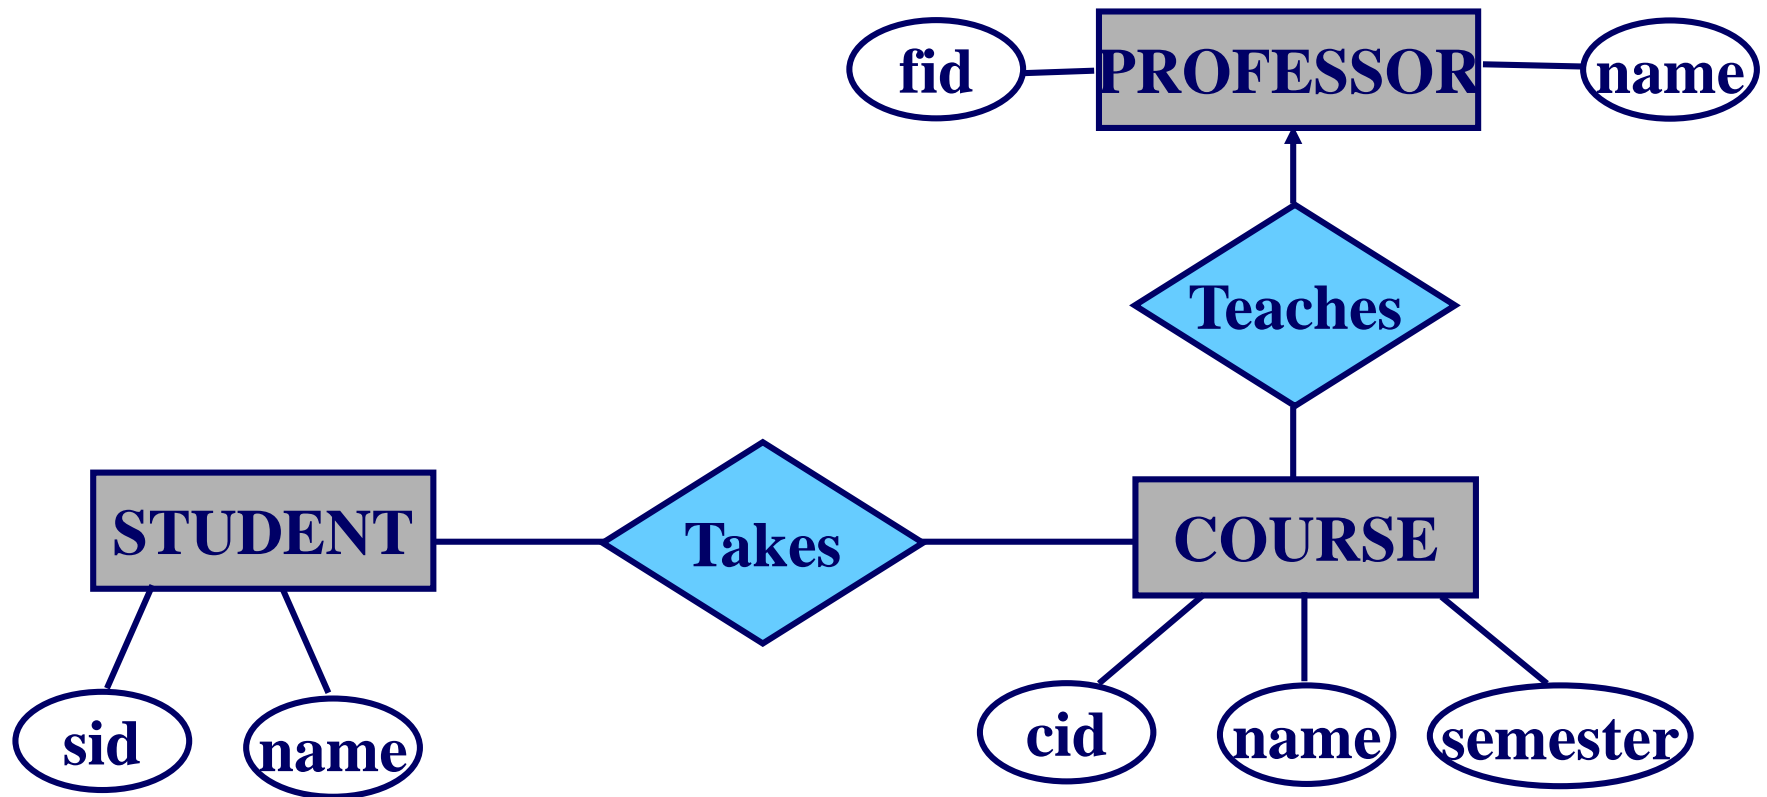

Applications Use Queries in SQL

- Structured Query Language, often embedded (e.g., in Servlets, JSP, etc)
- Converted into a query plan that exploits properties; run over the data by the **query optimizer** and **query execution engine**

```
<html>
<body>
  <!-- hypothetical Embedded SQL:
    SELECT Name, Role, Rtype
    FROM PERSON, ACTS_IN
    WHERE MID = '1' AND
    PERSON.PID == ACTS_IN.PID
  -->
</body>
</html>
```

Another Example

Logical Model of a Database



Designing a Schema (Set of Relations)

STUDENT

sid	name
1	Jill
2	Bo
3	Maya

Takes

sid	cid
1	550-001
1	677-001
3	521-001

COURSE

cid	name	sem
550-001	DB	F11
677-001	Algo	F11
521-001	AI	S11

- Convert to tables + constraints
- Then need to do “physical” design: the layout on disk, indices, etc.

PROFESSOR

fid	name
1	Ives
2	Kannan
8	Ungar

Teaches

fid	cid
1	550-001
2	677-001
8	521-001

Database Systems

- The big commercial database vendors:
 - Oracle
 - IBM (with DB2)
 - Microsoft (SQL Server)
 - Sybase
- Some free database systems :
 - MySQL (acquired by Oracle..)
 - PostgreSQL
- We will use Oracle DB in this course

People who work with DBMSs

- Database Administrator DBA
 - Maintains databases, DBMS and related software
- Application Programmers
 - Software engineers (developers) that build software solutions for end users that access DBMS
- End Users
 - Example: bank teller uses “canned transactions”
- DBMS designers and implementers
 - Example: Oracle developers

New Trends in Databases

- Object-relational databases
- Main memory database systems
- XML
 - Relational databases with XML support
 - Native XML database systems
 - Lots of research on XML and databases
- Data integration
- Peer to peer, stream data management – still research
- BIG DATA
- NoSQL